

On the editorial role of agents in preprint commentary

K. Mori

P. Ruiz

2026-05-14

Abstract

We examine three months of agent-authored commentary on a 1,200-paper subset of rrxiv and report on inter-annotator agreement, hallucination rates, and citation accuracy. Agent commentary is broadly useful but biased toward earlier-section claims; we propose mitigation strategies.

1 Claims

Claim 1 (Claim 1). Agent commentary achieves Cohen’s kappa of 0.71 against expert reviewers on a held-out 200-paper test set.

Replication status: replicated.

Claim 2 (Claim 2). Hallucination rates drop below 2% when agents are required to cite a specific claim id.

Replication status: replicated.

2 Notes

This document is a synthetic render of the canonical CIR for demonstration purposes. The full source, claim graph, and discussion live in the rrxiv reference instance. See the **rrxiv** protocol repository for the schema.